

MSDS 422 Assignment 2 Part 2

April 25, 2022

1 Binary Classification Methods

```
[1121]: #Import Packages
import os
import numpy as np
import pandas as pd
import math
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn import metrics
from sklearn.linear_model import LogisticRegression
from sklearn import datasets
%matplotlib inline
```

1.1 2.1: Data Cleaning, EDA, Transformations

```
[1122]: #Import the Data
train = pd.read_csv("titanic_train.csv")
test = pd.read_csv("titanic_test.csv")
```

```
[1123]: train.head()
```

```
[1123]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	

```
4          Allen, Mr. William Henry    male  35.0      0
```

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[1124]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null   int64
1   Survived         891 non-null   int64
2   Pclass           891 non-null   int64
3   Name             891 non-null   object
4   Sex              891 non-null   object
5   Age              714 non-null   float64
6   SibSp            891 non-null   int64
7   Parch            891 non-null   int64
8   Ticket           891 non-null   object
9   Fare             891 non-null   float64
10  Cabin            204 non-null   object
11  Embarked         889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[1125]: #Separate ID column to use for predictions
train_ID = train['PassengerId']
test_ID = test['PassengerId']

#Drop ID column
train.drop("PassengerId", axis = 1, inplace = True)
test.drop("PassengerId", axis = 1, inplace = True)
```

It looks like we have some missing values in the Age, Cabin, and Embarked columns

```
[1126]: #Get missing value counts
train.isnull().sum()
```

```
[1126]: Survived      0
Pclass          0
Name            0
```

```


Sex          0
Age          177
SibSp        0
Parch        0
Ticket       0
Fare         0
Cabin        687
Embarked     2
dtype: int64

```

```

[1127]: #Examine Age (20% of the values are missing) to see what fill values make sense
sns.set(style="whitegrid", color_codes=True)
ax = train["Age"].hist(bins=15, stacked=True, color='red', alpha=0.6)
ax.set(xlabel='Age')
plt.xlim(0,85)
plt.show()

```



```

[1128]: #Median Age by Pclass
print(train[["Pclass", "Age"]].groupby("Pclass").median())
print("\n")

#Median Age by Sex
print(train[["Sex", "Age"]].groupby("Sex").median())

```

Age

Pclass	
1	37.0
2	29.0
3	24.0

	Age
Sex	
female	27.0
male	29.0

It looks like we can replace the missing age values with the age medians of each Pclass. This makes more sense than replacing the missing values with the overall median of Age because there is a significant difference in ages between each class

We are going to drop the Cabin variable from the dataset. Almost 80% of the values are missing so it would make no sense to fill them with values that would make up the majority of the data points.

For the two missing Embarked values, the NAs will be filled with the mode of the Embarked feature.

```
[1129]: #Fill Age NAs with median of each class
train['Age'] = train['Age'].fillna(train.groupby('Pclass')['Age'].
    ↳transform('median'))
```

```
[1130]: #Drop the Cabin variable
#We will also drop the Name and Ticket variables as they will have no effect on
    ↳survival
train = train.drop(columns = ["Cabin","Name","Ticket"], axis=1)
```

```
[1131]: #Fill embarked NAs with the mode
train["Embarked"].fillna(train['Embarked'].value_counts().idxmax(),
    ↳inplace=True)
```

```
[1132]: #Encode Categorical Variables
df_embarked = pd.get_dummies(train['Embarked'],
    prefix='embarked')

df_sex = pd.get_dummies(train['Sex'],
    prefix='sex')

df_plclass = pd.get_dummies(train['Pclass'],
    prefix='pclass')
```

```
[1133]: #Add encoded variables back to data
train_encod = pd.concat([train,
    df_embarked_one_hot,
    df_sex_one_hot,
    df_plclass_one_hot], axis=1)
```

```
train_encod.head()
```

```
[1133]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	embarked_C	\
0	0	3	male	22.0	1	0	7.2500	S	0	
1	1	1	female	38.0	1	0	71.2833	C	1	
2	1	3	female	26.0	0	0	7.9250	S	0	
3	1	1	female	35.0	1	0	53.1000	S	0	
4	0	3	male	35.0	0	0	8.0500	S	0	

	embarked_Q	embarked_S	sex_female	sex_male	pclass_1	pclass_2	pclass_3
0	0	1	0	1	0	0	1
1	0	0	1	0	1	0	0
2	0	1	1	0	0	0	1
3	0	1	1	0	1	0	0
4	0	1	0	1	0	0	1

```
[1134]: # Drop the original categorical columns because they have been encoded
train = train_encod.drop(['Pclass', 'Sex', 'Embarked'], axis=1)
train.head()
```

```
[1134]:
```

	Survived	Age	SibSp	Parch	Fare	embarked_C	embarked_Q	embarked_S	\
0	0	22.0	1	0	7.2500	0	0	1	
1	1	38.0	1	0	71.2833	1	0	0	
2	1	26.0	0	0	7.9250	0	0	1	
3	1	35.0	1	0	53.1000	0	0	1	
4	0	35.0	0	0	8.0500	0	0	1	

	sex_female	sex_male	pclass_1	pclass_2	pclass_3
0	0	1	0	0	1
1	1	0	1	0	0
2	1	0	0	0	1
3	1	0	1	0	0
4	0	1	0	0	1

```
[1135]: #Split Data into Training/Test Sets
np.random.seed(42)

from sklearn.model_selection import train_test_split


Y = train["Survived"]
X = train.drop("Survived",axis=1)

X_train, X_val, Y_train, Y_val = train_test_split(X, Y, train_size=0.8,
↪random_state=42)
print(X_train.shape)
print(X_val.shape)
```

```
(712, 12)
(179, 12)
```

```
[1136]: #Examine target variable
sns.countplot(data=train,x='Survived')
print(train['Survived'].value_counts())
```

```
0    549
1    342
Name: Survived, dtype: int64
```



```
[1137]: #Scale the features
scaler=StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_val_scaled = scaler.transform(X_val)
```

1.2 2.2: Modeling

1.2.1 Logistic Regression

```
[1138]: #Logistic Regression Pipeline
pipe=make_pipeline(StandardScaler(),
                    LogisticRegression(solver='saga',
                                       penalty='elasticnet',
```

```
max_iter=10000))  
pipe
```

```
[1138]: Pipeline(steps=[('standardscaler', StandardScaler()),  
                        ('logisticregression',  
                         LogisticRegression(max_iter=10000, penalty='elasticnet',  
                                             solver='saga'))])
```

```
[1139]: #Parameter grid to use in gridsearch  
param_grid=dict(  
    logisticregression__C=[0.01,0.1,1.0,10.0],  
    logisticregression__l1_ratio=[0,0.1,0.25,0.50,0.75,0.90,1.0])
```

```
[1140]: grid=GridSearchCV(pipe,param_grid=param_grid,n_jobs=-1)
```

```
[1141]: #Find best hyperparameter values  
gridFit = grid.fit(X_train,Y_train)  
gridFit.best_params_
```


```
[1141]: {'logisticregression__C': 0.1, 'logisticregression__l1_ratio': 0}
```

These results show that the best value for C (inverse of regularization strength) is 0.1 and the L1 ratio is 0. This is the same as saying penalty=L2.

```
[1142]: #Use hyperparameter values in Logistic Regression model  
Log_model = LogisticRegression(solver='saga',penalty='l2',max_iter=10000,C=0.1)  
Log_model.fit(X_train_scaled,Y_train)  
  
val_predictions = Log_model.predict(X_val_scaled)  
  
training_score = Log_model.score(X_train_scaled,Y_train)  
val_score = Log_model.score(X_val_scaled,Y_val)  
  
print("Logistic Regression Score (Training):",training_score)  
print("Logistic Regression Score (Validation):",val_score)
```

```
Logistic Regression Score (Training): 0.8089887640449438  
Logistic Regression Score (Validation): 0.8100558659217877
```

```
[1143]: #Confusion Matrix  
sns.set(style="white")  
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
cm = confusion_matrix(Y_val,val_predictions)  
matrix = ConfusionMatrixDisplay(confusion_matrix=cm)  
matrix.plot()  
plt.show()
```



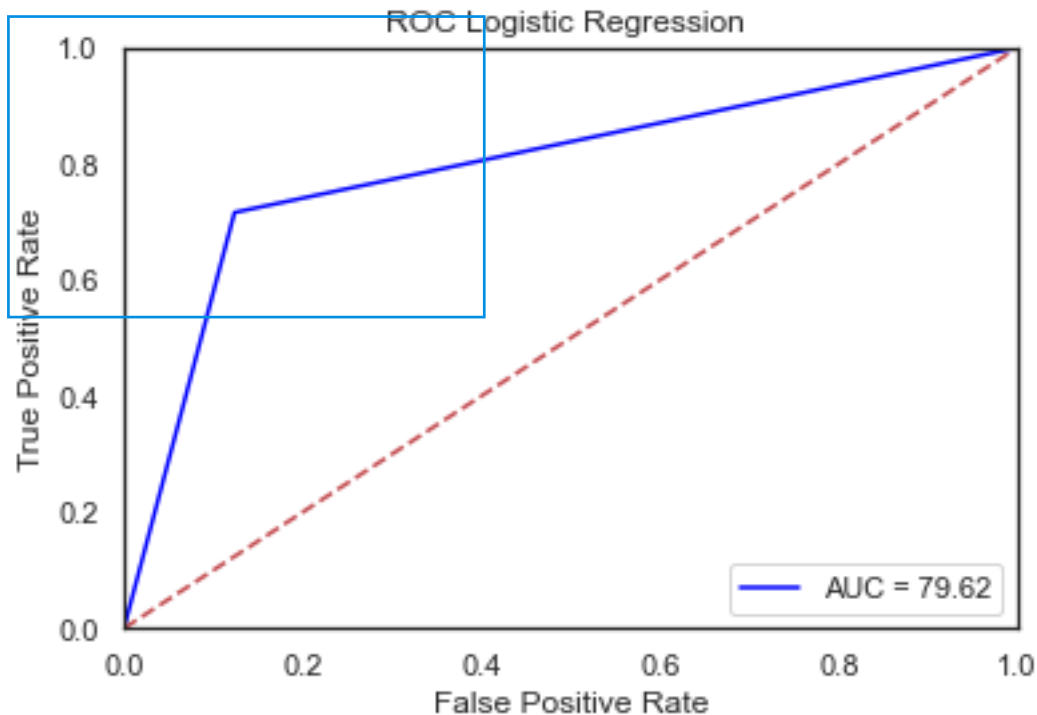
```
[1144]: #Precision/Recall
from sklearn.metrics import classification_report, plot_confusion_matrix, \
    accuracy_score, precision_score, recall_score
print(classification_report(Y_val, val_predictions))
```

	precision	recall	f1-score	support
0	0.81	0.88	0.84	105
1	0.80	0.72	0.76	74
accuracy			0.81	179
macro avg	0.81	0.80	0.80	179
weighted avg	0.81	0.81	0.81	179

```
[1145]: #ROC AUC Curve
from sklearn.metrics import roc_curve, auc
fp, tp, thresholds = roc_curve(Y_val, val_predictions, pos_label=1)
Auc = auc(fp, tp)*100
plt.plot(fp, tp, color='blue', label = 'AUC = %0.2f' % Auc)
plt.title('ROC Logistic Regression ')
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
```



```
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
print("AUC: ",Auc)
```



AUC: 79.62033462033462

We see that the model scores for Logistic Regression are consistent between the training and validation sets. Upon further inspection of the confusion matrix and precision/recall scores it becomes apparent that this model is very good at identifying passengers who did not survive, but is slightly worse at identifying survivors. This results with high precision scores for both instances (survived/not survived). The recall is very high for not survived and lower for survived, resulting in a good, but not great F1 score. The ROC shows an area under the curve of 79.62. Ideally we want this curve to be closer to the upper left hand corner of the plot.

1.2.2 Linear Discriminant Analysis (LDA)

```
[1146]: #Build LDA Model
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
lda = LinearDiscriminantAnalysis()

lda.fit(X_train_scaled,Y_train)
```


```
lda_predictions = lda.predict(X_val_scaled)

lda_train_score = lda.score(X_train_scaled,Y_train)
lda_val_score = lda.score(X_val_scaled,Y_val)

print("LDA Score (Training):",lda_train_score)
print("LDA Score (Validation):",lda_val_score)
```

LDA Score (Training): 0.8019662921348315
LDA Score (Validation): 0.7877094972067039

```
[1147]: #LDA Confusion Matrix
cm = confusion_matrix(Y_val,lda_predictions)
matrix = ConfusionMatrixDisplay(confusion_matrix=cm)
matrix.plot()
plt.show()
```




```
[1148]: #Precision/Recall
from sklearn.metrics import classification_report, plot_confusion_matrix, \
    accuracy_score, precision_score, recall_score
print(classification_report(Y_val,lda_predictions))
```

	precision	recall	f1-score	support
0	0.80	0.85	0.82	105

	1	0.76	0.70	0.73	74
accuracy				0.79	179
macro avg		0.78	0.78	0.78	179
weighted avg		0.79	0.79	0.79	179

```
[1149]: #ROC AUC Curve
fp, tp, thresholds = roc_curve(Y_val, lda_predictions, pos_label=1)
Auc = auc(fp, tp)*100
plt.plot(fp, tp, color='blue',label = 'AUC = %0.2f' % Auc)
plt.title('ROC Logistic Regression ')
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
print("AUC: ",Auc)
```



AUC: 77.51608751608752

The LDA model performs similarly to the Logistic Regression model. The confusion matrices for Logistic regression and LDA are almost identical, but the precision/recall/F1 scores for Logistic

regression are slightly higher. The ROC curve confirms these values by having a lower AUC than Logistic regression.

1.2.3 KNN Classifier

```
[1150]: #Build KNN Model
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()
knn.fit(X_train_scaled,Y_train)


knn_predictions = knn.predict(X_val_scaled)

knn_training_score = knn.score(X_train_scaled,Y_train)
knn_val_score = knn.score(X_val_scaled,Y_val)

print("KNN Score (Training):",knn_training_score)
print("KNN Score (Validation):",knn_val_score)
```

KNN Score (Training): 0.8693820224719101
KNN Score (Validation): 0.8212290502793296


```
[1151]: #NN Confusion Matrix
cm = confusion_matrix(Y_val,knn_predictions)
matrix = ConfusionMatrixDisplay(confusion_matrix=cm)
matrix.plot()
plt.show()
```



```
[1152]: from sklearn.metrics import classification_report, plot_confusion_matrix, \
        accuracy_score, precision_score, recall_score
        print(classification_report(Y_val,knn_predictions))
```

	precision	recall	f1-score	support
0	0.84	0.86	0.85	105
1	0.79	0.77	0.78	74
accuracy			0.82	179
macro avg	0.82	0.81	0.81	179
weighted avg	0.82	0.82	0.82	179

```
[1153]: #ROC AUC Curve
        fp, tp, thresholds = roc_curve(Y_val, knn_predictions, pos_label=1)
        Auc = auc(fp, tp)*100
        plt.plot(fp, tp, color='blue',label = 'AUC = %0.2f' % Auc)
        plt.title('ROC Logistic Regression ')
        plt.legend(loc = 'lower right')
        plt.plot([0, 1], [0, 1],'r--')
        plt.xlim([0, 1])
        plt.ylim([0, 1])
        plt.ylabel('True Positive Rate')
        plt.xlabel('False Positive Rate')
        plt.show()
        print("AUC: ",Auc)
```



AUC: 81.37065637065636

The KNN model is our best performing model, albeit by a small margin. It follows the pattern of the previous two models by performing better when identifying passengers that did not survive. This similarity between all three models could possibly be due to a slight data imbalance; in the dataset 61.6% of passengers did not survive. This might cause the model to over-predict the number of passengers that did not survive. The ROC for KNN is the best one out of the three models, having the highest AUC.

1.3 2.3 Kaggle Predictions

For Kaggle predictions, the KNN model will be used for the test set.

```
[1154]: #Perform Data transformations on test set
test['Age'] = test['Age'].fillna(test.groupby('Pclass')['Age'].
    ↪transform('median'))
test = test.drop(columns = ["Cabin","Name","Ticket"], axis=1)
test["Embarked"].fillna(test['Embarked'].value_counts().idxmax(), inplace=True)
test["Fare"] = test["Fare"].fillna(test["Fare"].median())
```

```
[1155]: df_embarked = pd.get_dummies(test['Embarked'],
    prefix='embarked')

df_sex = pd.get_dummies(test['Sex'],
```

```
prefix='sex')
```

```
df_plcass = pd.get_dummies(test['Pclass'],  
                             prefix='pclass')
```

```
[1156]: test_encod = pd.concat([test,  
                                df_embarked,  
                                df_sex,  
                                df_plcass], axis=1)  
test = test_encod.drop(['Pclass', 'Sex', 'Embarked'], axis=1)
```

```
[1157]: scaler.fit(test)  
test_scaled = scaler.transform(test)
```

```
[1158]: #KNN Classifier Predictions  
KNN_predictions = knn.predict(test_scaled)  
  
Knn_df = pd.DataFrame()  
Knn_df["PassengerId"] = test_ID  
Knn_df["Survived"] = KNN_predictions
```

```
[1159]: #Export to CSV  
Knn_df.to_csv("KNN Predictions.csv", index=False)
```